

SEAL: Statistically-Bounded Erasure Audits for Machine Unlearning under Mutual Distrust

Quang-Vinh Dang

Abstract—A data owner who submits an erasure request to a machine learning system operator has no direct way to verify that the request was honored, and the operator has no incentive to volunteer proof of non-compliance. Prior work shows that any audit accurate enough to catch a dishonest operator necessarily leaks information about the retained population, but stops short of a mechanism a mutually distrusting pair can actually run. This paper presents SEAL, a two-channel erasure-audit certificate that separates a zero-privacy-cost check on the data owner's own forgotten records from a differential-privacy-bounded check on synthetic boundary probes, calibrated against replicated honest re-executions rather than assumed thresholds, and its counterpart SEAL-W for large language model unlearning, which retrieves a randomly chosen member of a paraphrased query cluster through real Paillier-based private information retrieval and cross-checks response diversity against a phantom-topic baseline. We state five formal results (four theorems and a proposition) – mostly standard tools (a two-hybrid IND-CPA reduction, a hypergeometric evasion bound, an exact Student- t predictive threshold, Boole's inequality) applied to this certificate's specific channels, not new mathematics – precisely scoping what each channel provably defeats: private information retrieval obliviousness under the decisional composite residuosity assumption, a hypergeometric bound on a resource-limited adversary's evasion probability, an exact finite-sample calibrated threshold, a zero-diversity guarantee against uniform response suppression, and a Bonferroni-corrected composition bound whose violation we found in our own implementation, and fixed, by writing this proof down. We reproduce the verification methodology of WaterDrum, the closest prior watermark-based unlearning audit, as an explicit baseline: in one of three trials at this paper's smallest calibration-sample size, a server that detects and filters the audit query itself defeats the WaterDrum-equivalent single-query check – the exact gap WaterDrum's own authors report as unsolved – while SEAL-W's diversity channel catches it; the other two trials illustrate the same small-sample calibration limitation this paper reports throughout, rather than a clean three-for-three win. Experiments on federated logistic regression over UCI German Credit, MNIST, and a 150,000-row Forest Covertype subsample, and on a real fine-tuned language model audited over the real TOFU fictitious-author benchmark, report both power and honestly-measured failure modes at small calibration-sample sizes, rather than tuning them away.

Index Terms—Machine unlearning, differential privacy, private information retrieval, text watermarking, auditing, federated learning, right to be forgotten.

1 INTRODUCTION

Data-protection regulation gives individuals a right to have their data erased from systems trained on it [1]. When

the system is a machine-learned model, “erasure” is a claim about a training procedure the data owner cannot observe: the model operator asserts that a retraining or fine-tuning step removed the record's influence, and the data owner must decide whether to believe it. This is not a benign estimation problem. The operator may find honest unlearning expensive and have every incentive to claim compliance while changing nothing (*lazy noncompliance*), or to narrowly patch only the behavior an audit is expected to probe while leaving the record's broader influence untouched (*targeted spoofing*).

Two recent results sharpen why this is hard rather than merely inconvenient. Tang, Joshi, and Kundu give an information-theoretic argument that, for convex models, any audit relying only on the model's behavior (query-response pairs) cannot separate an honestly-unlearned model from an insufficiently-unlearned one without revealing membership information about the retained set [2] – there is an unavoidable trade-off between how powerful a behavioral audit is and how much it leaks about records the data owner was never entitled to see. Nguyen et al.'s veriFUL framework independently surveys the federated-unlearning-verification landscape and concludes that no proposed scheme yet gives a federation of mutually distrusting clients and servers a concrete, run-once verification protocol with a stated privacy cost [3]. Both papers characterize the problem; neither gives an auditor a mechanism to run.

1.1 Contributions

This paper's contribution is a concrete audit protocol – not a new impossibility result, and not a claim to have invented differential privacy, private information retrieval (PIR), or text watermarking. It is the specific combination of known primitives into a certificate that (i) makes the audit power/privacy-leakage trade-off in [2] a dial the data owner turns rather than an unavoidable fact they simply accept, and (ii) instantiates that dial concretely for two settings veriFUL leaves open: federated classical-ML unlearning, and large-language-model unlearning verified via watermarking. Concretely:

- 1) A two-channel certificate (SEAL) that separates a *zero-marginal-privacy-cost* channel – the data owner re-evaluates the model on data it already possesses in full, which costs no retained-set privacy

Q.-V. Dang is with British University Vietnam, Hung Yen, Vietnam (e-mail: vinh.dq4@buv.edu.vn). ORCID: 0000-0002-3877-8024.

at all – from a *tunable-privacy-cost* channel using Gaussian-mechanism-noised synthetic boundary probes, whose spend is set by a stated (ϵ, δ) budget (Section 3.2).

- 2) A calibration procedure that builds each channel’s null distribution from literal re-executions of the honest mechanism, matched to the real run’s base seed and per-client ordering, rather than an assumed parametric null – a detail we show empirically is not optional (Section 3.2).
- 3) SEAL-W, a watermark-based certificate for LLM unlearning that plants several independently-paraphrased queries about the same forgotten entity in a committed batch, retrieves one at random via real Paillier-based PIR, and adds an independent response-diversity channel targeting exactly the adversary class – uniform, topic-level output filtering – that a single-query watermark check structurally cannot catch (Section 3.3).
- 4) Five formal results (four theorems and a proposition), precisely scoping – not inventing – what each channel is proven to defeat (Section 4.1; full derivations in the supplemental material). Four combine standard tools (a two-hybrid IND-CPA reduction, the hypergeometric distribution, an exact Student-*t* predictive interval via Cochran’s theorem, Boole’s inequality) with this certificate’s specific channels, rather than proving new mathematics; their value is precision about scope, not proof difficulty. The fifth, the Bonferroni composition bound, found something we did not already know before proving it: a violation present in our own first implementation, caught by writing the proof down rather than by an experiment failing visibly, and fixed prior to the experiments reported here.
- 5) An explicit reproduction of WaterDrum’s verification methodology [4] as a first-class experimental baseline, and a demonstration that it is defeated by exactly the server behavior its own authors report as an open problem, while SEAL-W’s diversity channel is not (Section 4.2).
- 6) An empirical evaluation (Section 5) on three classical-ML datasets spanning three orders of magnitude in per-client data volume, and on a real fine-tuned language model audited over the real TOFU fictitious-author benchmark [5], reporting failure modes – including a benchmark-scale finding that audit power *degrades* as per-client data grows – as honestly as the successes.

Every individual primitive reused here – the Gaussian mechanism and zero-concentrated differential privacy (zCDP) accounting [6], Paillier-based PIR [7], the Kirchenbauer et al. green-list watermark [8], and the general oblivious-audit-via-PIR pattern of RESPIR [9] – is standard and cited rather than re-derived, per the guidance that theoretical derivations already published elsewhere belong in the reference list, not repeated. What is new is their specific composition into a certificate that answers a stated open problem in the unlearning-verification literature, and the formal accounting of exactly which adversary class each

piece of that composition defeats.

2 RELATED WORK

2.1 Unlearning Verification under Mutual Distrust

Machine unlearning itself – the underlying capability SEAL audits, rather than provides – has an extensive literature independent of any audit mechanism, starting with Cao and Yang’s original formulation of unlearning as a summation-form retraining problem [10]. SISA training [11] restructures training into isolated, retrainable shards so that forgetting one point requires retraining only its shard, trading a constant-factor storage and computation overhead for an exact retraining guarantee; Guo et al.’s certified removal [12] and Sekhari et al.’s follow-up [13] instead give approximate, statistically certified removal guarantees for convex (or convex-adjacent) models via a bounded parameter perturbation, avoiding retraining entirely but only for the model classes and update budgets their perturbation analysis covers. None of these three papers addresses verification: each assumes the party performing the deletion is honest and asks only how to delete efficiently, not how a second, distrusting party can check that deletion occurred. The federated setting has its own unlearning-method literature with the same gap – Liu et al.’s FedEraser [14] accelerates client-level retraining by reusing cached historical updates, again assuming the server that runs it is honest – and Gao et al.’s VeriFi [15] and Sommer et al.’s Athena [16] are the closest existing verification-side proposals: VeriFi unifies several federated-unlearning and verification-method combinations into one empirical framework for measuring unlearning’s effect, while Athena gives a client a backdoor-membership-inference hypothesis test to run against a claimed-forgotten model. Both, like WaterDrum (below), assume a non-adaptive or honest-but-curious server rather than the actively adversarial one SEAL’s threat model covers (Section 3.1), and neither states a controllable privacy cost for the verification check itself – Table 1 records this gap alongside the audit-specific comparisons that follow.

Tang, Joshi, and Kundu’s geometric transfer theorem [2] is, to our knowledge, the sharpest existing formal statement of why this problem is hard: for convex models, it derives a per-point lower bound relating a behavioral audit’s accuracy to the leakage it must incur on a named retained point, and shows this bound is essentially tight for the class of audits it covers. It is a *characterization* result – it says what any sufficiently powerful behavioral audit must cost, not how to build one that spends that cost deliberately, chooses where to spend it, or composes multiple audit channels with different leakage profiles. Nguyen et al.’s veriFUL survey [3] independently reaches the same practical conclusion from the federated-learning side: after cataloguing verification entities, goals, and metrics proposed across the federated-unlearning literature, it concludes that none of the surveyed schemes gives a federation of mutually distrusting clients and servers a concrete, run-once protocol with a stated, controllable privacy cost. Neither paper is in tension with SEAL’s design; SEAL is best read as an attempt at the mechanism both papers imply is missing, evaluated against the specific adversary classes (lazy noncompliance, targeted

spoofing) that motivate their impossibility and survey arguments respectively.

González et al.’s sequential auditing framework [17] solves an adjacent but structurally different problem: given repeated, i.i.d. access to a mechanism under two auditor-chosen adjacent datasets, it gives an anytime-valid test of whether the mechanism satisfies a *claimed* (ϵ, δ) -DP guarantee, with dramatically fewer samples than earlier batch tests. This does not transfer directly to the erasure-audit setting: the data owner here does not control a repeatable mechanism it can re-run on chosen adjacent inputs (it gets one trained model), and there is no claimed (ϵ, δ) for the *training* process to audit against – the quantity SEAL bounds is the leakage of its own *audit protocol*, not a property of the server’s unknown, unaudited training pipeline. Table 1 summarizes how these three lines of work, plus Godinot et al.’s RESPIR (below), relate to SEAL/SEAL-W along the dimensions that matter for the mutual-distrust setting.

2.2 Watermark-Based Data Provenance

Kirchenbauer et al. introduce a keyed, context-seeded green/red vocabulary partition for LLM generation, detected via a z -test on the green-token fraction [8]; Lau et al.’s Waterfall generalizes this to a vocabulary-permutation-and-perturbation scheme applied via an LLM paraphraser, targeting robustness and scalability for large-scale text provenance rather than unlearning specifically [18]. Lu et al.’s WaterDrum applies this family of watermarks specifically to unlearning verification, and is the closest prior work to SEAL-W [4]: its data-centric metric watermarks a data owner’s contribution and checks for the watermark’s continued presence after a claimed unlearning step. Its own appendix explicitly states that a server which detects and filters the watermark check itself is not handled by their construction – Section 4.2 reproduces this gap directly, as an experimental comparison against our own implementation of their verification methodology, rather than only citing the limitation as stated.

2.3 Oblivious Auditing via Private Information Retrieval

Godinot et al.’s RESPIR uses lattice-based (LWE) PIR to let an auditor commit to a batch of candidate queries and retrieve a checked subset obliviously, with planted canaries as tamper sentinels, for fairness audits of a model provider that may adapt its answers once it can distinguish an audit query from a genuine one [9]. SEAL-W adopts this committed-batch-with-canaries pattern – the auditor commits the server to an answer for every candidate before revealing which will be checked – but substitutes Paillier-based, single-server PIR [7] for RESPIR’s lattice-based construction. This is a deliberate simplification, not an oversight: lattice-based PIR is designed for sublinear communication at large database sizes, which matters for RESPIR’s fairness-audit setting (potentially large candidate pools of real users) but not at the batch sizes used here (tens of candidates per audit round), where Paillier’s linear server-side work and one-ciphertext-per-index communication are entirely practical and rest on a simpler, more standard hardness assumption (Section 4.1, Theorem 1).

2.4 Differential Privacy Accounting and Auditing

Channel B’s privacy cost is accounted via zero-concentrated differential privacy (zCDP) [6], chosen for its clean, tight composition over the B independent probe queries within one audit round. The Gaussian differential privacy / f -DP framework of Dong, Roth, and Su [19] gives the general Neyman–Pearson-optimal trade-off curve between a hypothesis test’s power and its leakage for a mechanism whose noise distribution is known exactly to the auditor – of which SEAL’s actual setting is a strictly harder special case, since the auditor knows only a finite calibration sample standing in for that distribution, not the distribution itself. We do not claim to solve, or even fully characterize, this harder finite-calibration variant of the f -DP tradeoff in general; Section 6.2 states this gap as an open question rather than papering over it with an unproven claim of optimality. Channel B’s leakage measurement (Section 5) is itself an instance of the empirical DP-auditing-via-membership-inference paradigm that Jagielski et al. [20] and Nasr et al. [21] developed to measure what a differentially private mechanism’s implementation actually leaks, as opposed to its accounted worst-case bound: both give the adversary a released statistic (there, a poisoned or canary training point; here, the noised loss-gap on synthetic boundary probes) and measure how well it distinguishes membership empirically, exactly the audit-side counterpart to the accounting-side guarantee Channel B spends against.

2.5 Influence-Based Spoofing Detection

Channel B’s ability to catch narrow, targeted fine-tuning (rather than genuine retraining) that a zero-cost own-data check cannot rests on the same intuition as Koh and Liang’s influence functions [22]: honest retraining moves the loss on a forgotten record’s *neighborhood* in feature space, not only the exact record itself, because the record’s removal changes the model’s fit in a region, not at a point; a narrowly targeted fine-tune that memorizes only the exact forgotten records does not reproduce this neighborhood effect. SEAL’s boundary-probe construction (perturbing D_f into nearby synthetic points, Section 3.2) operationalizes this intuition as a statistical test rather than as an influence-function computation, which would itself require access to the model’s internals the data owner does not have.

3 SYSTEM DESIGN

3.1 Threat Model and Problem Formulation

3.1.1 Setting and Notation

A data owner (or a regulator acting on its behalf) submitted a request to erase a record set D_f from a model trained by an operator (“the server”) on $D = D_f \cup D_r$, where D_r is the retained set. The server returns a model M' (parameters θ' ; the model before the request had parameters θ) and claims D_f ’s influence has been removed. The data owner does not trust this claim, cannot inspect M' ’s weights, training log, or optimizer state, and interacts with M' only through whatever query interface the deployment already exposes (label/logit queries for the classical track; text generation for the LLM track). The data owner possesses D_f in full (it is, definitionally, the data whose erasure it requested) but

TABLE 1
Positioning SEAL/SEAL-W against the closest prior work.

Work	Setting	Runnable protocol?	Stated privacy budget?	Mutual distrust?	Crypto. oblivious?
Tang et al. [2]	Behavioral unlearning audit	No (impossibility bound)	Characterized, not chosen	Yes	N/A
Nguyen et al. [3]	Federated unlearning (survey)	No (identifies the gap)	Not addressed	Yes	N/A
González et al. [17]	Auditing a claimed DP guarantee	Yes, for repeatable mechanisms	Type-I error, not audit leakage	No (chosen datasets)	N/A
Godinot et al. [9]	Fairness audits, deceptive model	Yes	Not privacy-budgeted	Yes	Yes (lattice-based)
VeriFi [15]	Federated unlearning verification	Yes	Not addressed	Partially (honest-but-curious)	No
Athena [16]	Backdoor-MI unlearning verification	Yes	Not addressed	No (client-side test only)	No
WaterDrum [4]	LLM unlearning verification	Yes, single canonical query	Not addressed	Partially (non-adaptive server)	No
SEAL / SEAL-W	Both settings above	Yes	Yes, tunable (ϵ, δ)	Yes	Yes (Paillier, SEAL-W)

TABLE 2
Notation used throughout the paper.

D_f, D_r θ, θ'	Forgotten / retained record sets, $D = D_f \cup D_r$ Model parameters before / after the claimed unlearning step
$l(\theta, x, y)$	Per-sample loss
Λ_A, Λ_B	SEAL Channel A / Channel B audit statistics
(ϵ, δ)	Differential-privacy budget spent by Channel B
$\sigma, L_{\max}, B$	Channel B noise scale, loss-clip bound, probe count
R (classical)	Number of calibration replicates (Section 3.2)
R (SEAL-W)	Real-cluster / phantom-cluster size (Section 3.3)
m, N, P $\beta_{\text{target}}, \beta_{\text{ch}}$	Tampered / total / checked slots (Theorem 2) Target and per-channel (Bonferroni-split) false-accusation rate
$\tau, q(\beta, m)$	Calibrated threshold, exact predictive quantile (Theorem 3)
z, z_{i^*}	Watermark z -score; the one PIR-retrieved real-cluster score

does *not* possess, and is not entitled to learn anything further about, D_r . Table 2 collects the symbols used throughout the rest of the paper.

3.1.2 Adversary Classes

Table 3 lists the server behaviors in scope, in increasing order of sophistication. Each is a strategy for producing M' (and, for SEAL-W, an accompanying response policy) given full knowledge of the audit protocol's construction – a Kerckhoffs's-principle-style assumption: only the calibration data and the auditor's specific random draws in a given round are secret, not the protocol's design.

Two adversary classes are explicitly *out* of scope, stated here rather than left implicit: a server that adapts its training or filtering strategy *after* observing which specific audit round it is being tested in (SEAL-W's PIR obliviousness, Theorem 1, is precisely what rules this out within a round, but a server that learns something from the aggregate pattern of many audit rounds over time is not analyzed here); and a server with a costless, perfect semantic classifier willing to refuse *all* queries – audit and genuine alike – about an entity, discussed further in Section 6.1.

TABLE 3
Adversary classes in scope, and which channel targets each.

Behavior	Description	Targeted by
Honest	Genuinely retrains/fine-tunes excluding D_f	(the null hypothesis every channel is calibrated against)
Lazy	Ignores the request; keeps training on all of D, D_f included	Channel A / SEAL-W slot channel
Targeted spoof	Fine-tunes only on D_f (classical track) to fool an anticipated own-data check, leaving D_r -wide influence untouched	Channel B (boundary probes detect the missing neighborhood effect)
Filter-only (LLM)	Output-filters the one exact known canonical audit string	SEAL-W real-cluster (any paraphrase defeats it)
Topic-filter (LLM)	Output-filters <i>any</i> phrasing mentioning the forgotten entity, uniformly	SEAL-W diversity channel (Prop. 1)
Partial topic-filter (LLM)	Output-filters a guessed subset of m known paraphrasings, resource-bounded	SEAL-W slot channel (Thm. 2)

3.1.3 Design Goal

The data owner's goal is a decision rule T , built only from D_f , a disjoint calibration corpus, and query access to M' , that:

- 1) is calibrated to a stated false-accusation rate β : $\Pr[T = \text{dishonest} \mid \text{honest}] \leq \beta$;
- 2) carries a stated, and where possible zero, bound on how much running T reveals about D_r – the quantity Tang et al.'s theorem [2] shows cannot be driven to zero *in general* for a sufficiently powerful behavioral test, but which SEAL and SEAL-W make an explicit, chosen dial rather than an unacknowledged side effect.

3.2 SEAL: Auditing Classical Federated Unlearning

3.2.1 Two-Channel Certificate

Let θ, θ' be the model's parameters before and after the claimed unlearning step, and $l(\theta, x, y)$ the per-sample loss.

Channel A (zero cost). The data owner computes $\Lambda_A = \text{mean}_{(x,y) \in D_f} [l(\theta', x, y) - l(\theta, x, y)]$ entirely on data it already possesses. No noise is added and no privacy budget is spent, because the released quantity is a function only of the requester's own data and the public model.

Channel B (tunable cost). The data owner perturbs its own forgotten records into synthetic boundary probes $x_i + \epsilon_j$, $\epsilon_j \sim \mathcal{N}(0, \text{jitter}^2 I)$, and asks the server to evaluate M' on them, label held fixed. Each returned per-probe loss is clipped to $[0, L_{\max}]$ and released through a Gaussian mechanism at noise level σ , giving Λ_B analogous to Λ_A but over the noised probe losses. Because the probes are synthetic points constructed by the auditor from its own data – not real retained records – the Gaussian mechanism's sensitivity bound applies directly, and zCDP composition over B probes gives $\rho_{\text{total}} = BL_{\max}^2/(2\sigma^2)$, converted to (ϵ, δ) via $\epsilon = \rho_{\text{total}} + 2\sqrt{\rho_{\text{total}} \ln(1/\delta)}$ [6]. Given a target (ϵ, δ) , this inverts to the noise level σ the server's responses must carry – turning the impossibility of [2] into a dial: choosing ϵ fixes exactly how much retained-set privacy this audit round may cost.

3.2.2 Calibration

Both channels are compared against a null distribution built by literally re-executing the honest procedure R times, on a disjoint public calibration corpus standing in for D_f , using the *same base random seed and per-client ordering* as the real training run. This detail is not cosmetic: warm-starting a calibration replicate from an already-converged model, or varying the base seed across replicates, under-integrates or adds spurious trajectory variance to the null and produces a degenerate, near-zero-variance threshold in our own early implementation attempts (Section 4.1 gives the corrected, exact threshold this null feeds into).

3.2.3 Certificate Algorithm

Algorithm 1 summarizes the full data-owner-side procedure. The server-side interface is minimal and identical across every honest or dishonest mechanism in Table 3: given a batch of inputs, return M' 's per-sample losses. This is deliberate – the certificate's power comes entirely from which inputs are queried and how the responses are combined, not from any cooperation the server must specially provide, so an honest server's implementation burden is exactly the burden of answering ordinary inference queries.

3.3 SEAL-W: Auditing LLM Unlearning

A data owner contributed watermarked text (via [8], keyed per contributor) that was used to fine-tune a language model, and wants to check a claimed erasure of its influence.

3.3.1 Why a Cluster, Not One Query

A server recognizing one exact known trigger string is defeated by any paraphrase. A smarter server can instead recognize the *topic* (the forgotten entity's name, however worded) and either suppress every phrasing uniformly, or hardcode canned replies for only a guessed subset of phrasings it can afford. Neither a single paraphrased query nor PIR alone defeats either variant: against uniform suppression there is no “which slot is checked” secret worth

Algorithm 1 SEAL certificate (data-owner side)

Require: D_f , model access to θ (before) and θ' (after), calibration corpus D_c , target β_{target} , DP budget (ϵ, δ) , jitter scale, B probes, R calibration replicates

- 1: $\Lambda_A \leftarrow \text{mean}_{(x,y) \in D_f} [l(\theta', x, y) - l(\theta, x, y)]$ {Channel A, zero cost}
- 2: $\sigma \leftarrow \text{CALIBRATE_SIGMA}(\epsilon, \delta, B, L_{\max})$ {zCDP inversion, [6]}
- 3: Draw probes $\{x_i + \epsilon_j\}_{j=1}^B$, $\epsilon_j \sim \mathcal{N}(0, \text{jitter}^2 I)$
- 4: Query server for $\{l(\theta', x_i + \epsilon_j, y_i)\}$, clip to $[0, L_{\max}]$
- 5: $\Lambda_B \leftarrow \text{mean}_j[\text{noisy_}l(\theta', \text{probe}_j) - l(\theta, \text{probe}_j)]$ {Channel B}
- 6: **for** $r = 1$ to R **do**
- 7: Retrain fresh from θ' 's base seed on D_c -substituted-for- D_f shard, same client ordering
- 8: Honestly fine-tune to exclude the substituted shard
- 9: Record $(\Lambda_A^{(r)}, \Lambda_B^{(r)})$ on this replicate
- 10: **end for**
- 11: $(\mu_A, \sigma_A), (\mu_B, \sigma_B) \leftarrow \text{mean/std of the } R \text{ replicate values}$
- 12: $\beta_{\text{ch}} \leftarrow \beta_{\text{target}}/2$ {Bonferroni split, Thm. 4}
- 13: $\tau_A \leftarrow \mu_A - q(\beta_{\text{ch}}, R) \sigma_A$; $\tau_B \leftarrow \mu_B - q(\beta_{\text{ch}}, R) \sigma_B$ {Thm. 3}
- 14: **return** **dishonest** if $\Lambda_A < \tau_A$ or $\Lambda_B < \tau_B$, else **honest**, with (ϵ, δ) spent and β_{target} controlled

hiding, since the server treats every phrasing alike; against the resource-bounded guesser, hiding which slot is checked is exactly the right game, but requires more than one real query to hide among.

3.3.2 Protocol

Commit. The auditor assembles a candidate batch: an R -member *real cluster* of independently-paraphrased queries about the forgotten entity, one positive-control canary query about a known-retained entity, an R -member *negative-canary cluster* of paraphrased queries about a phantom entity never present in any training data, and decoy queries about other retained entities. The server answers every query – it cannot yet tell which will be checked. Each answer's watermark z -score [8] and a 32-bit SimHash response fingerprint [23] are quantized to integers and committed via a SHA-256 hash over both arrays, so the server cannot revise an answer once it learns what matters.

Retrieve. The auditor retrieves, via real single-server Paillier PIR [7]: one uniformly-random real-cluster member's z -score (chosen only now, not at commit time), the positive canary, every negative-canary-cluster member (for the diversity baseline), and every real-cluster fingerprint.

Decide. Two independent channels, plus canary sanity checks. The *slot channel* flags if the one retrieved real-cluster z -score exceeds a calibrated upper threshold. The *diversity channel* flags if the real cluster's mean pairwise Hamming diversity, relative to the phantom cluster's, falls below a calibrated lower threshold. Both thresholds, and the two canary sanity bounds, are calibrated from held-out calibration authors exactly as in Section 3.2, with the false-accusation budget β_{target} split across all active conditions (Theorem 4). Algorithm 2 summarizes the full auditor-side procedure.

Algorithm 2 SEAL-W certificate (auditor side)

Require: forgotten entity f , retained entities, phantom entity p , real-cluster size R , n_{decoys} , PIR keypair, calibration authors, β_{target}

- 1: **Commit** (server side): build batch = R paraphrases of $f \cup 1$ positive canary $\cup R$ paraphrases of $p \cup n_{\text{decoys}}$ decoys; shuffle
- 2: Server answers every prompt in the batch (cannot yet tell which will be checked)
- 3: For each answer: compute watermark z -score [8] and 32-bit SimHash fingerprint [23]; quantize both to integers
- 4: Commit = SHA-256(database || fingerprints)
- 5: **Retrieve** (auditor side, real Paillier PIR [7]): $i^* \sim \text{Uniform}(\{1, \dots, R\})$ real-cluster index; retrieve z_{i^*} , positive-canary z , all R negative-canary z 's and fingerprints, all R real-cluster fingerprints
- 6: **Decide:** build nulls $\{z^{(r)}, \text{div-ratio}^{(r)}, z_{\text{pos}}^{(r)}, z_{\text{neg}}^{(r)}\}_{r=1}^{R_{\text{calib}}}$ by re-running honest unlearning targeting each calibration author instead of f
- 7: $n \leftarrow$ number of conditions calibrated in this call (up to 4); $\beta_{\text{ch}} \leftarrow \beta_{\text{target}}/n$ (Thm. 4)
- 8: $\tau_{\text{slot}}, \tau_{\text{div}}, \tau_{\text{pos}}, \tau_{\text{neg}} \leftarrow$ predictive quantiles at β_{ch} from each null (Thm. 3)
- 9: **return** **dishonest** if $z_{i^*} > \tau_{\text{slot}}$ or $\text{div-ratio} < \tau_{\text{div}}$ or $z_{\text{pos}} < \tau_{\text{pos}}$ or $z_{\text{neg}} > \tau_{\text{neg}}$, else **honest**

In the experiments of Section 5, $R = 5$ (the full paraphrase-template pool, so the hypergeometric prediction of Theorem 2 matches the resource-bounded adversary's guessed subset one-for-one), $n_{\text{decoys}} = 2$, and $R_{\text{calib}} = 2$ calibration authors – the minimum for a defined sample variance, and, as Section 6.1 discusses at length, the binding constraint on this track's measured reliability.

3.3.3 Grounding the Author Pool in a Real Benchmark

The audited fictitious-author profiles can be built two ways. The default experiment invents author names and generates each author's watermarked training documents from scratch. A second, additive source instead grounds the same protocol in the real, published TOFU benchmark [5]: TOFU's "full" split lays out 200 fictitious authors as 4,000 question/answer rows in 200 contiguous 20-row blocks with no explicit author-id column, so real names and biographical facts must first be recovered from the raw text. We do this with two deterministic heuristics, hand-verified against three independently-known authors and swept across all 200 blocks: an author's name is the single most frequent two-word run of consecutive capitalized words across the block's text (TOFU answers restate the full name far more often than any incidental two-capitalized-word phrase, such as a book title), with a trailing possessive normalized so that "X" and "X's" mentions reinforce one tally rather than splitting it; a short fact snippet is the first sentence of the first answer matching a fixed keyword-priority order (genre, award, profession, notable work, birth date), skipping bare name restatements. Because the name heuristic is a plurality vote over bigram counts rather than a verified extraction, we additionally measured its confidence margin – the gap between the winning bigram's count and the runner-up's

– across all 200 blocks: the typical block is unambiguous (median count 35, median margin 27 over the runner-up, no block with zero extracted candidates), but a real minority is not – 12 of 200 blocks (6%) end in an exact tie between two equally-frequent bigrams, and 30 of 200 (15%) have a margin of 2 or fewer, i.e. a plausible alternative candidate (a co-author, a recurring book title) nearly as frequent as the winner. We report this honestly as the heuristic's actual error-rate exposure rather than only the three-author spot-check: a non-trivial fraction of the 200-block pool may seed watermarked generation with a wrong or ambiguous name, though this affects only which real name/fact seeds a synthetic training document (Section 3.3.3's scope note), not any certificate mechanism, threshold, or theorem.

Critically, this does *not* mean training on TOFU's literal question/answer text: that text is static, pre-existing, and carries no watermark bias, so using it directly as a training document would produce a document with no green-list signal at all and silently defeat the entire certificate. Instead, the recovered name and fact snippet only seed the same, unmodified watermarked-generation call ([8], Section 3.3) that the default experiment uses with an invented name – every training document is still our own watermarked generation, now about a real entity and a real fact rather than a fabricated one. The reserved negative-canary pool is drawn from the same real TOFU distribution and held out of all fine-tuning entirely across a run, rather than hand-invented, for the same reason: a genuine, distributionally-matched "never trained on" control is a strictly better negative canary than an author name chosen by the auditor with no textual grounding at all. This substitution changes only where the author pool's content comes from, not the certificate, the mechanisms, or any threshold in Sections 3.2–4.1; it is not comparable to the official TOFU leaderboard, which fine-tunes and evaluates phi-1.5 or LLaMA-2-7B against TOFU's own forget-quality metrics, rather than `distilgpt2` against a watermark/PIR certificate.

4 THEORETICAL GUARANTEES AND COMPARISON TO PRIOR ART

4.1 Theoretical Guarantees

We state five results precisely; complete proofs are in the supplemental material. Each states exactly, and only, what it establishes.

Theorem 1 (PIR obliviousness). *If Paillier is IND-CPA secure (equivalently, the Decisional Composite Residuosity assumption holds), then for any two indices $i \neq j$ in an n -slot database, the one-hot-encrypted PIR query for i is computationally indistinguishable from that for j , with $\text{Adv}_{\text{index-prio}}(A) \leq 2 \text{Adv}_{\text{IND-CPA}}(B)$ for an adversary B built from A via a two-hybrid argument.*

Theorem 2 (Hypergeometric slot-hiding bound). *For a batch of N slots, P of which the auditor checks (uniformly at random, independent of the adversary's choice), and an adversary that tampers with m slots before learning which will be checked,*

$$\Pr[\text{checked slots} \cap \text{tampers slots} = \emptyset] = \binom{N-m}{P} / \binom{N}{P}.$$

Theorem 3 (Exact finite-sample calibrated threshold). *Let $X_1, \dots, X_m \sim \mathcal{N}(\mu, \sigma^2)$ i.i.d. (calibration replicates) and $Y \sim \mathcal{N}(\mu, \sigma^2)$ independent (the real audit statistic under the null). Then $(Y - \bar{X})/(S\sqrt{1+1/m}) \sim t_{m-1}$, so $\tau = \bar{X} + t_{m-1, 1-\beta} S\sqrt{1+1/m}$ satisfies $\Pr[Y > \tau] = \beta$ exactly. This threshold strictly exceeds the asymptotic z -quantile plug-in at every finite m , and converges to it as $m \rightarrow \infty$.*

Proposition 1 (Exact zero diversity under uniform suppression). *If a server answers every query mentioning a topic with one fixed string, then for any real-cluster retrieval of $R \geq 2$ such queries, the mean pairwise SimHash Hamming diversity is exactly 0.*

Theorem 4 (Bonferroni composition, and the bug it caught). *If a certificate’s dishonesty decision is the logical OR of k independently-calibrated conditions each controlled at level β_i , the overall false-accusation rate is bounded by $\sum_i \beta_i$ (Boole’s inequality), so a target rate β_{target} requires $\beta_i = \beta_{\text{target}}/k$. Our first SEAL-W implementation calibrated all $k=4$ conditions (slot, diversity, two canary checks) at the same, un-split β_{target} , bounding the true rate only at $4\beta_{\text{target}}$ – a gap found by writing this theorem down and checking the code against it, not by an experiment failing visibly, and fixed prior to the experiments reported here.*

4.2 Reproducing and Breaking WaterDrum

WaterDrum’s verification methodology – watermark a contribution, then test whether post-unlearning generations still carry the watermark – is, architecturally, exactly SEAL-W’s slot channel run in isolation: one canonical (unparaphrased) query, one watermark score, one calibrated threshold, no diversity channel, no real-cluster. We implement this literal reduction as an explicit baseline (`waterdrum_baseline_decide`) and run it side by side with the full SEAL-W decision on the same fine-tuned model and the same four server behaviors, honestly noting one substitution: WaterDrum’s own implementation embeds its watermark via Waterfall’s paraphraser-and-permutation scheme [18], while we use the simpler Kirchenbauer green-list watermark [8] throughout, for tractability on the CPU-only hardware used in our experiments. This is a real substitution of one statistical detection mechanism for another, not a like-for-like reproduction, and we do not claim the two watermarks have identical detection power, false-positive behavior, or robustness to the server’s own filtering under all conditions – only Waterfall’s own published evaluation can speak to Waterfall’s specific numbers. What the substitution does not change is the structural claim this experiment is built to isolate: the failure mode demonstrated below is that a server which detects and filters the audit query defeats a single-canonical-query check regardless of which watermark scored that query, because the filter acts before the watermark score is ever computed. A reader who doubts this generalizes across watermark families should read Table 4 as evidence about the single-query-versus-diversity-channel architecture, not as a direct benchmark of Kirchenbauer’s watermark against Waterfall’s.

Table 4 reports the result against a server that uniformly suppresses every query mentioning the forgotten entity’s name (`topic_filter`). Every trial shows the real cluster’s

TABLE 4
WaterDrum-equivalent baseline vs. full SEAL-W against `topic_filter`, per trial ($N_{\text{calib}} = 2$, $\beta_{\text{target}} = 0.10$).

Trial	Real-cluster div.	SEAL-W	WaterDrum-baseline
0	0.0 (thr. -1.02)	No	Yes (also flags honest)
1	0.0 (thr. -0.38)	No	No
2	0.0 (thr. $+0.46$)	Yes	No

response diversity collapsing to *exactly* zero (Proposition 1 proves this must happen, not merely that it tends to), and in the one trial where the calibration sample was not itself too degenerate to support any confident call, SEAL-W’s diversity channel flags the attack while the WaterDrum-equivalent single-query check does not – reading a filtered, non-watermarked canned reply as indistinguishable from genuine forgetting. This is precisely the false negative WaterDrum’s Appendix D reports as unsolved, reproduced directly against the mechanism built to exploit it, not merely cited.

Trials 0 and 1 illustrate a second, independent finding: at $N_{\text{calib}} = 2$, the Bonferroni-corrected threshold (Theorem 4) can itself land below zero, so even the true, exact-zero diversity ratio this adversary always produces is not always low enough to cross it. This is not a defect specific to WaterDrum-style checks; it is the same finite-calibration-sample limitation that governs every threshold in this paper, surfacing here in the one comparison built specifically to expose it.

5 EXPERIMENTAL EVALUATION

5.1 Classical Track

Federated multinomial-logistic/softmax models are trained across 5 simulated clients via FedAvg-style local-then-average aggregation [24] on three datasets of increasing scale: UCI German Credit [25] (a bank-consortium right-to-be-forgotten instantiation, ~ 40 -record forget sets), MNIST [26] (5,000-sample subset, 784-dim), and a 150,000-row sub-sample of UCI Forest Covertype [27] (54-dim, $\sim 30,000$ rows/client). Table 5 reports detection power against *lazy* noncompliance (power at which SEAL’s zero-cost Channel A alone separates honest from lazy) against two baselines: raw parameter distance, and an unrestricted behavioral audit in the style of [2]. This figure is the AUC of the relevant raw statistic across honest-vs-dishonest Monte Carlo trials – an oracle-threshold measure of the statistic’s separability, computed identically for SEAL and both baselines so the three-way comparison is apples-to-apples – and is deliberately distinct from the calibrated certificate’s actual single-instance pass/fail decision rate under Theorem 3’s exact finite-sample threshold. That calibrated flag rate is reported separately (Table 7) and is substantially lower at this dataset’s production calibration budget ($n_{\text{calib}} = 25$): 0.225, not 0.920. Both numbers are true and neither supersedes the other – the AUC measures the statistic’s raw information content independent of calibration cost, while the calibrated flag rate measures what a real audit round, with a finite calibration budget, actually outputs – but we report both explicitly, rather than let Table 5’s AUC stand in unqualified for the certificate’s realized single-instance behavior.

TABLE 5

Detection power vs. lazy noncompliance (Channel A alone, zero privacy cost) vs. two baselines. This is the raw statistic's AUC, not the calibrated certificate's flag rate – see Table 7 and prose.

Dataset	Param.-diff	Unrestr. behav.	SEAL (Ch. A)
German Credit	0.743	0.553	0.920
MNIST	0.597	0.500 (chance)	1.000
Covertypes (150k)	0.610	0.510	0.540

At higher input dimension (MNIST), the unrestricted behavioral baseline performs at exact chance: broad queries over a 784-dimensional retained population dilute one client's ~ 50 forgotten records into noise, exactly the diffuseness [2]'s theorem predicts costs power. SEAL's zero-cost channel needs no such broad query and reaches 1.000.

Table 6 reports the designed power/privacy-leakage tradeoff against *targeted spoofing* (fine-tuning only on D_f), across a sweep of Channel B's privacy budget ε , together with retained-set leakage measured as the empirical membership-inference attack AUC an adversary achieves given exactly what Channel B releases. Concretely, the "attacker" here is not a separately trained classifier but the released statistic itself, read as a score: the (possibly noised) loss-gap $l(\theta', x, y) - l(\theta, x, y)$ is computed on a held-out pool of same-distribution, never-trained-on records (the "non-members") and on genuinely retained records (the "members"), and the AUC of separating the two using exactly that one released number *is* the leakage measure – an honestly worst-case-for-the-defender proxy, since it hands the adversary the released quantity directly rather than requiring it to be learned from limited access. SEAL's power climbs monotonically with ε toward the unbounded-noise limit, while leakage stays close to chance (0.50–0.54) across the whole sweep on German Credit and MNIST – with the label-preserving synthetic probes used here, the near-neighbor membership signal is intrinsically weak at these datasets' scale (small forget sets, ~ 40 – 50 records), so a tight ε does not have to suppress much signal in the first place. Parameter-diff also catches this specific spoof at zero accounted privacy cost, but only because the spoof mechanism used here is an aggressive, large, narrowly-directed parameter update; a subtler spoof would not show up in parameter distance at all, and there is no way to know which regime a real deployment is in without an audit like Channel B's.

Read honestly, Table 6 also states a real limitation of Channel B's current probe design, not only a success: at every genuinely privacy-bounded budget tested ($\varepsilon \leq 4$), power against this spoof stays in the 0.525–0.653 range – barely above the 0.50 chance line – and only reaches 1.000 at $\varepsilon = 10^6$, which the table itself labels "no noise," i.e. not a private setting at all. Framing Channel B as "a dial the data owner turns" is accurate about the mechanism (Theorem 1's and the zCDP composition's $\varepsilon \rightarrow \sigma$ inversion is exact and does hand the data owner a chosen operating point), but it should not be read as a claim that every point on that dial is a *useful* one against this spoof: the useful end of the dial, for the probe construction used here, is close to the unbounded-noise limit. Whether a different boundary-

TABLE 6

SEAL Channel B: power (raw-statistic AUC, not the calibrated flag rate) vs. targeted spoofing and retained-set leakage (attack AUC), across privacy budget ε .

ε	German Credit	MNIST	Leakage (both)
0.05	0.525	0.542	0.50–0.53
0.50	0.528	0.556	0.50–0.53
1.0	0.538	0.556	0.50–0.53
4.0	0.583	0.653	0.50–0.53
10^6 (no noise)	1.000	1.000	0.50–0.53

probe design (larger perturbation magnitude, a different loss statistic, or more probes per audit round) moves the useful regime to a smaller, genuinely private ε is a concrete open question this evaluation raises but does not answer.

Figure 1 plots detection power against retained-set leakage directly for all three datasets and both baselines, making the same comparison visually: SEAL's own-data channel (against lazy noncompliance) sits at zero measured leakage cost by construction, strictly above both baselines at every leakage level they occupy, on German Credit and MNIST; against targeted spoofing, SEAL's curve sweeps from near-chance power at tight ε to the two baselines' fixed operating point as ε relaxes, tracing out the tradeoff Table 6 reports numerically.

Covertypes' result is reported as a genuine, honest limitation, not tuned away: at $\sim 30,000$ rows/client, SEAL's own-data channel loses most of its power (0.540, barely above the 0.500 chance line and the two baselines), because a federation with far more data per client is intrinsically more robust to losing any one client's contribution – the audit signal "did removing this data change how well the model fits it" is weakest exactly where the model has the most other data to fall back on. This is visible directly in Figure 1's bottom panel: unlike the top two datasets, SEAL's lazy-detection point is barely separated from the baselines rather than sitting in a clearly dominant region.

The certificate's calibrated single-instance decision (Theorem 3's exact threshold) is more conservative than the raw statistic: German Credit's false-accusation rate fell from 0.15 to 0.10 and MNIST's from 0.083 to 0.0 after replacing the asymptotic z -quantile plug-in with the exact predictive quantile, with no change to the Monte-Carlo power numbers above (those do not depend on the threshold formula).

To test this paper's own repeated claim that additional calibration replicates, not a different threshold formula, is the fix for small-sample instability (Section 6.1), we ran a dedicated ablation on German Credit at $\varepsilon = 1$, holding every other hyperparameter at Table 10's value and varying only n_{calib} across $\{3, 5, 10, 15, 20, 25\}$, at $n_{\text{trials}} = 80$ (rather than this dataset's usual 20) specifically so that outer Monte-Carlo sampling noise (resolution $1/n_{\text{trials}}$ per estimated rate) stays well below the effect size being measured – an initial run at $n_{\text{trials}} = 15$ produced non-monotonic rates indistinguishable from pure sampling noise. Table 7 reports the result honestly, including where it complicates the paper's own narrative: the certified detection rate against lazy non-compliance does show a real, if noisy, improving trend with more replicates ($0.113 \rightarrow 0.225$ from $n_{\text{calib}} = 3$ to 25), consistent with more calibration data producing a tighter, more




Fig. 1. Detection power (raw-statistic AUC; see Table 7 for the calibrated certificate’s actual flag rate) vs. retained-set leakage (attack AUC), SEAL’s ϵ -sweep against both baselines: German Credit (left), MNIST (center), Covertypes (right). Each panel’s left subplot is against lazy noncompliance, right subplot against targeted spoofing.

TABLE 7

German Credit calibration-replicate ablation ($\epsilon=1$, $n_{\text{trials}}=80$): certified rates as n_{calib} varies, all other hyperparameters fixed at Table 10’s values.

n_{calib}	False-acc. rate	Detect. (lazy)	Detect. (spooft)
3	0.025	0.113	0.025
5	0.088	0.250	0.038
10	0.025	0.188	0.013
15	0.025	0.188	0.000
20	0.025	0.213	0.013
25	0.025	0.225	0.013

informative threshold; but the false-accusation rate does *not* show a clean monotonic improvement – it is flat at 0.025 for five of the six sampled values and is, if anything, worse at $n_{\text{calib}} = 5$ (0.088) than at $n_{\text{calib}} = 3$. This is consistent with what Section 6.1 argues on separate grounds: at small n_{calib} , which specific replicates happen to be drawn dominates the resulting threshold, so the false-accusation rate at any one small value of n_{calib} is itself a noisy quantity rather than a smooth function of n_{calib} , and this six-point sweep is not powered to distinguish a genuine non-monotonicity in that noisy quantity from ordinary sampling variation across the independent draws used at each value. The raw, uncalibrated Channel A statistic’s power against lazy noncompliance is unaffected by n_{calib} by construction (0.879 throughout, matching Table 5’s German Credit entry within Monte Carlo noise); every number in Table 7 instead reflects only the calibrated certificate’s pass/fail decision, not the underlying signal’s raw strength.

5.2 LLM Track

distilgpt2 is fine-tuned (CPU-only) on watermarked biographies of fictitious authors drawn from the real TOFU benchmark [5], audited by SEAL-W against four server behaviors of increasing sophistication (*lazy*, exact-string *filter-only*, uniform *topic-filter*, and a resource-bounded *partial topic-filter* that guesses 50% of a 5-member paraphrase pool). Table 8 tracks four successive, honestly-measured threshold fixes on the same experiment.

Each fix moves both numbers in the direction the corresponding theorem predicts *before* it was run: an exact threshold cannot manufacture confidence a small calibration sample does not support, so tightening the threshold to be exact, and then to be correctly composed across channels, necessarily costs detection power as it removes false accusations. These rates should be read for that qualitative direction, not as precise point estimates: with only 3 trials, a rate of 1/3 carries a 95% Clopper–Pearson interval of

TABLE 8

Successive certificate fixes, same experiment (3 trials, $N_{\text{calib}}=2$; each rate’s 95% Clopper–Pearson interval spans roughly [0.01, 0.91] – see prose).

Version	False-acc. rate	Detect. topic-filt.
Hardcoded canary floor	1.0 (degenerate)	1.0
Calibrated, asymptotic z	0.667	1.0
Calibrated, exact Student- t	0.333	0.667
+ Bonferroni split (Thm. 4)	0.0	0.333

TABLE 9

Full SEAL-W certificate, final (Bonferroni-corrected) version, all four server behaviors (3 trials, $N_{\text{calib}}=2$, $\beta_{\text{target}}=0.10$; rates here carry the same wide 3-trial confidence intervals as Table 8). The last row is a theoretical prediction, not an empirical rate – see prose.

Server behavior	Detection rate (empirical)
Honest (false-accusation rate)	0.0
Lazy (keeps training on everything)	0.333
Filter-only (exact-string canned reply)	0.667
Topic-filter (uniform suppression)	0.333
– diversity channel alone	0.333
Partial topic-filter (guesses 50% of $R=5$)	0.333
– analytic evasion <i>probability</i> (theory)	0.6

roughly [0.01, 0.91], wide enough that “consistent with” a theorem’s prediction means the observed count does not contradict it, not that the 3-trial experiment independently corroborates the prediction with any real precision. The large- n_{trials} German Credit ablation (Table 7, Section 6.1) is this paper’s higher-confidence evidence for the calibration-replicate-count claim these small- N LLM-track tables can only illustrate qualitatively, at the sample sizes a CPU-only, real-fine-tuned-model experiment makes tractable.

Table 9 reports the final, Bonferroni-corrected certificate’s detection rate against all four server behaviors, together with the diversity channel’s rate in isolation and the resource-bounded adversary’s analytic evasion prediction. The partial-topic-filter adversary guesses $m = 2$ of the $R = 5$ paraphrase templates (a guess fraction of 0.5, rounded half-to-even by Python’s `round`, so $0.5 \times 5 = 2.5$ rounds to the even integer 2, not 3); the resulting empirical detection rate (0.333) is consistent with the exact analytic hypergeometric prediction (Theorem 2: $1 - (R - m)/R = 1 - 3/5 = 0.4$) within this trial count’s sampling noise – the one place in this evaluation where a component is checked against a closed-form formula rather than only Monte Carlo. The WaterDrum-comparison result (Section 4.2, Table 4) is drawn from this same final run.

This 3-trial, $N_{\text{calib}}=2$ regime is small by design: each trial

TABLE 10

Classical-track (SEAL) hyperparameters, exact per dataset, sufficient to reproduce every number in Tables 5–6.

Parameter	German Credit	MNIST	Coverttype
Clients	5	5	5
FedAvg rounds (train)	25	20	20
Unlearn extra rounds	8	6	6
Local epochs, lr	1, 0.3	1, 0.3	1, 0.3
Batch size	32	64	512
Calibration replicates R	25	15	12
Probe count B , jitter	50, 0.4	50, 0.05	50, 0.3
Monte-Carlo trials	20	12	10
β_{target}	0.05	0.05	0.05

fine-tunes and audits a real language model on CPU-only hardware, and every rate above should be read as “roughly this many out of 3,” not a precise estimate – the qualitative pattern (which mechanism is caught by which channel, and that both fixes moved every rate in the theoretically-predicted direction) is the result this scale supports, not the exact percentages. Section 6.1 returns to this as the paper’s central open lever.

5.3 Implementation and Reproducibility

Every mechanism, channel, and adversary discussed above is implemented, not simulated: the federated softmax model and both classical-track channels are pure-Python/NumPy; the LLM track fine-tunes a real `distilgpt2` checkpoint via PyTorch/Hugging Face Transformers; the watermark is the published green-list construction [8], not a proxy for it; and SEAL-W’s PIR is a real Paillier implementation (the `phe` library), with the auditor’s queries genuinely encrypted end-to-end rather than the obliviousness property being asserted without a corresponding cryptographic artifact. All four adversary mechanisms (lazy, filter-only, topic-filter, partial-topic-filter) share identical underlying model weights within each experiment – a single “kept training on everything, never actually unlearned” fine-tune – differing only in the output-filtering wrapper placed over it, so that the comparison isolates what each certificate channel can detect rather than accidentally comparing different models. The full implementation, experiment scripts, and raw per-trial result files are available at <https://github.com/vinhqdang/SEAL-Statistically-bounded-Erasure-Audit-under-mutual-distrust>; the paper’s numbered theorem statements are additionally restated with complete, checkable proofs in the supplemental material. Tables 10–11 list every hyperparameter needed to reproduce the numbers in this section exactly, including per-dataset values that differ across the classical-track sweep.

A note on dual use, engaging directly with an objection worth taking seriously rather than dismissing: Table 3’s adversary taxonomy and Section 6.1’s explicit gaps (Section 6.1’s “What No Channel in This Paper Rules Out”) are, read one way, a roadmap – a self-interested operator who wants to evade this specific certificate now has a numbered list of the strategies known not to be caught (a standing, costless total topic ban; content-preserving noise injected into an otherwise-canned reply to defeat the diversity channel’s specific SimHash-based check) rather than having to

TABLE 11

LLM-track (SEAL-W) hyperparameters, sufficient to reproduce every number in Tables 8–9.

Base model	<code>distilgpt2</code> (82M params, CPU-only)
Fine-tune epochs, lr, doc repeats	$3, 5 \times 10^{-5}, 8$
Watermark γ (green fraction), δ (logit bias)	0.5, 8.0
Documents/author, generation length	4, 30 tokens
Authors/trial, calibration authors	5, 2
Real-cluster / phantom-cluster size R	5, 5
Decoys, Monte-Carlo trials	2, 3
β_{target}	0.10, split across up to 4 channels

discover them independently. We do not think this framing is wrong; we think it is incomplete in a way that still favors publishing. First, an operator sophisticated enough to run and defend a fielded unlearning pipeline is, in our experience, already capable of deriving these same evasion strategies from the public architecture of *any* behavioral audit – a total refusal to discuss a topic defeats a query-response test by definition, with or without this paper, and noise-plus-suppression is the generic countermeasure to any fixed similarity check, not a technique this paper originates. Second, disclosure is asymmetric in a way that matters here: this paper’s caveats are architecture-level (what *no* instantiation of this certificate family can rule out), while the certificate’s actual per-deployment decision boundary – the calibrated threshold of Theorem 3, computed from calibration replicates the auditor holds privately – is not published and cannot be reconstructed from this paper alone; an operator who knows the general evasion *category* still does not know the specific numeric line a live deployment will draw, or the exact calibration draws that determine it. This argument, however, does not cover every conceded gap equally, and we should say so rather than let it read as a blanket rebuttal: the total-topic-ban and content-preserving-noise-injection strategies conceded above are both *threshold-independent* – an operator executing either does not need to know the calibrated decision boundary at all, because the strategy defeats the channel structurally, not by landing just under a guessed threshold. Argument two therefore does not close either gap; it only means that a *subtler* evasion, one that tries to survive detection by a narrow margin rather than eliminate the signal outright, is harder to tune without the private threshold. Third, and least satisfying but most honest: of the two threshold-independent gaps, only the total-topic-ban strategy has an externally visible cost outside this paper’s audit – a model that provably, permanently refuses to discuss one specific named entity is itself a conspicuous, observable fact a user, journalist, or regulator can notice without running any audit protocol at all. The content-preserving-noise-injection gap has no such tell: a canned reply wrapped in innocuous paraphrasing or filler tokens looks, to anyone but the diversity channel’s specific SimHash-based check, like an ordinary response, and this paper does not have a countervailing argument for that case beyond stating it plainly as open (Section 6.1). We do not think a structurally similar field resolves this tension

by keeping its methods secret, for what it is worth: nuclear-material inspection regimes, the closest real-world analog to an audit run under mutual distrust between an inspector and a party with an incentive to cheat, publish their verification protocols and calibrate residual risk through inspection frequency and cost rather than through secrecy of method [28] – but that is a structural precedent for why disclosure does not forfeit deterrence in this class of problem, not a demonstration that this paper’s specific noise-injection gap is already closed. It is not; we are stating that directly rather than let the three preceding arguments imply otherwise. None of this makes the roadmap objection disappear; it is why we still judge disclosure the better call: an audit whose coverage is secret cannot be trusted by the mutually distrusting party it is meant to serve, and the alternative to stating these gaps precisely is not a more secure certificate, only a less honest description of the one that exists. The construction does not itself enable any new attack on a model or its training data; it is strictly a verification tool run by the data owner’s side of the protocol.

6 DISCUSSION, FUTURE WORK, AND CONCLUSION

6.1 Discussion and Limitations

6.1.1 Calibration Sample Size Is the Binding Constraint

The remaining gap between SEAL(-W)’s design and its measured single-instance reliability is governed by calibration-sample size, not by the threshold formula. Theorem 3’s correction is exact given the Gaussian idealization of the calibration replicates, but at $N_{\text{calib}} = 2$ (this paper’s LLM-track minimum, needed just to define a sample variance at all) which two replicates happen to be drawn dominates the resulting threshold more than the mechanism under test does – visible directly in Table 4’s per-trial spread, where the same true, exact-zero diversity signal is caught in one trial and missed in the other two purely because the calibrated threshold itself moved. Every fix applied in this paper (Table 8) made the certificate’s stated guarantee *more* truthful and, as a direct and expected consequence, less powerful at this sample size; none of them, and no other threshold formula, can substitute for more replicates. This is not a defect specific to the LLM track: the classical track’s Channel B shows the same pattern at its own, larger but still finite R (Section 3.2), and Coverttype’s result below shows a second, independent way finite information can bind. A dedicated calibration-replicate ablation on German Credit (Table 7) tests this claim directly rather than only arguing it: detection power against lazy noncompliance does improve, noisily but with a clear overall trend, as n_{calib} grows from 3 to 25 – but the false-accusation rate does not improve monotonically, and is worse at $n_{\text{calib}} = 5$ than at $n_{\text{calib}} = 3$. We report this honestly as a partial complication rather than a clean confirmation: more replicates helps power, as claimed, but the false-accusation rate at any single small n_{calib} is itself a high-variance quantity that a six-point, 80-trial sweep cannot fully separate from sampling noise, so “more replicates monotonically shrinks the false-accusation rate” should be read as the guarantee’s asymptotic direction, not as a claim this paper’s own finite ablation demonstrates cleanly at every sampled point.

6.1.2 Scale Is Not a Free Win

Coverttype’s power loss is a genuine, dataset-scale-dependent limitation of the own-data channel, not a universal one, and is reported as such rather than tuned away by choosing a friendlier configuration. The mechanism is intuitive in retrospect but was not assumed going in: a federation with $\sim 200\times$ more rows per client than German Credit converges to a substantially more stable optimum, so continuing to train on everything (the “lazy” adversary) perturbs the model’s fit on any one client’s forgotten records by much less, relative to ordinary round-to-round training noise at that scale, than it does with a few hundred rows per client. The audit signal Channel A relies on – “did removing this data change how well the model fits it” – is intrinsically weakest exactly where the model has the most other data to fall back on. Whether a tunable-cost channel analogous to Channel B, or a larger calibration budget, restores power at this scale is left as a direct, concrete next experiment rather than a claim made here.

6.1.3 What No Channel in This Paper Rules Out

Neither SEAL-W channel defeats a hypothetical server with a perfect, costless semantic classifier willing to refuse *all* queries – ordinary and audit alike – about an entity, i.e. a server behaviorally indistinguishable from a total, standing topic ban rather than a targeted, audit-specific spoof. No black-box behavioral test, with or without PIR-hidden retrieval, can formally rule this out: if the server is willing to make M' refuse to discuss an entity under any circumstance forever, no query-only audit can distinguish “never learned about this entity” from “knows about it and refuses to ever say so.” This is the same scope boundary Tang et al.’s impossibility framing draws for behavioral audits in general [2], and the same kind RESPIR’s own theorem statements draw for their PIR construction [9] – stating the covered adversary class precisely, rather than an unconditional claim, is the standard this literature already holds itself to.

A second, practical caveat concerns the diversity channel specifically: Proposition 1 is exact for literal, single-string uniform suppression, but a server that adds cheap, content-preserving noise to an otherwise-canned reply (paraphrasing its own filler text, injecting innocuous filler tokens) could in principle defeat a naive SimHash-based diversity check without restoring genuine informativeness, without being caught by the specific guarantee proved here. Studying this specific arms race – diversity-channel robustness against a server that knows it is being diversity-checked – is out of scope for this paper, which evaluates the concrete adversary it implements and measures (uniform canned-reply suppression), not every conceivable output-filtering strategy.

6.1.4 Legal Standing and Who Bears the Audit’s Cost

This paper motivates itself from GDPR’s Article 17 erasure right [1] and the CJEU’s Google Spain judgment [29], which established that a data controller can be compelled to act on such a request against its own interest, but neither settles a question this paper’s own construction raises directly: what evidentiary weight a calibrated statistical certificate, with a stated but nonzero false-accusation rate β , would actually carry with a data protection authority or a court.

No jurisprudence yet says whether a $\beta=0.05$ -controlled flag is treated as proof of noncompliance, as probable cause justifying a compelled audit under existing regulatory powers, or as inadmissible absent a stronger standard; this paper takes no position on that legal question and states plainly that it has not resolved it. A related, more practical point this paper’s own threat-model statement already half-concedes (Section 3.1: “a data owner, or a regulator acting on its behalf”): running SEAL-W in particular – real Paillier PIR, committed candidate batches, replicated calibration re-executions – is not something an individual data subject is realistically going to do by hand. In practice this is machinery for a well-resourced actor (a regulator, a consumer-advocacy body, or an aggregator running the audit on many data subjects’ behalf) rather than a tool each affected individual runs themselves, and the paper should not be read as claiming otherwise.

A second, separate question is who pays for the audit once someone does run it. Channel B is framed throughout as “a dial the data owner turns,” but the privacy budget that dial spends is not the data owner’s own: it is drawn from the retained-set records D_r – other people’s data, still in the model, who never requested this audit, were not asked to consent to it, and have no mechanism in this paper’s design to learn that their data’s leakage budget was spent, by how much, or how often. One data subject’s exercise of an erasure right thus imposes a real, measured privacy cost (Table 6’s leakage-AUC column is exactly this cost) on a different, non-consenting set of data subjects, unilaterally, each time the audit runs. This is a genuine distributive-fairness question this paper does not solve: nothing here bounds how many times D_r ’s privacy budget can be spent across repeated audit requests from different data owners, nor gives D_r ’s members any say in, or visibility into, that spending. We flag it explicitly rather than let the “dial” framing imply the cost is costless because it is chosen – chosen by the auditor, not by the party who pays it.

6.2 Future Work: An Open Theoretical Question

The certificate design in Sections 3.2–3.3 is engineering, in the sense used throughout this paper: a specific, checkable composition of known primitives, not a new theorem about the fundamental limits of calibration-based auditing. A genuinely open, and considerably harder, question is whether those limits admit a clean characterization at all. We state it here as a research agenda, not a result, and explicitly flag any future claim to have solved it in a single sitting as suspect until checked line by line.

The calibration-only erasure audit game. A server holds $D = D_f \cup D_r$ and an unlearning claim for D_f . The auditor does *not* know D_r (the entire point of mutual distrust) or the server’s mechanism family, and knows only D_f , a disjoint calibration corpus D_c of size m , and query access to M' . The auditor builds a test using only these, calibrating thresholds empirically from replicated honest runs on D_c – exactly what both certificates in this paper already do. Two quantities characterize the test jointly, as functions of the calibration budget m and query budget q : power $P(m, q)$ at a controlled false-accusation rate, and aggregate retained-set leakage $L(m, q)$, against a server that is not honest-but-curious but adversarially adaptive, choosing how to

respond to every query to best evade the test while knowing the test’s construction (though not its specific random draws).

This formalization is deliberately posed to specialize back to the known frameworks under extra assumptions: if the server is honest-but-curious and D_r reduces to a single known point, it should recover Tang et al.’s bound [2]; if the mechanism’s noise family is known exactly and $m \rightarrow \infty$, it should recover the f -DP tradeoff function [19]. Verifying these reductions is itself the first, necessary step – if the formalization does not specialize correctly, it is the wrong model and needs revision before anything built on it is worth trusting.

Three concrete conjectures organize an attack on this question, in increasing order of difficulty and decreasing order of how well existing tools apply: (A) *distribution-free achievability* – a calibration-only test achieving power within $O(1/\sqrt{m})$ of an oracle test that knows the mechanism’s noise distribution exactly, via standard concentration inequalities on the empirical calibration mean/variance; (B) *a matching minimax lower bound* against the adaptive, protocol-aware adversary above, via a Le Cam- or Fano-style argument where the two hypotheses being distinguished are two *server strategies* rather than two datasets – the least well-trodden of the three, and the one most likely to require a genuinely new argument; and (C) *adversarial multi-query composition*, treating the audit as a repeated game between an auditor allocating a continuous falsification-response budget and a no-regret learner, closer to online-learning regret bounds than to classical hypothesis testing. The most likely outcome of attempting (A) is that it reduces to a known concentration-inequality corollary once correctly stated, which would still be a real, if modest, contribution (the formalization itself); the most interesting possible outcome of (B) is a clean impossibility result with an explicit, non-obvious dependence on m . Either way, this is realistically months of focused work by someone with a research background in hypothesis testing or differential-privacy theory, and is offered here as a scoped, falsifiable agenda for that work, not a preview of its answer.

6.3 Conclusion

SEAL and SEAL-W give a mutually distrusting data owner and model operator a concrete, calibrated, formally-scoped audit certificate where prior work characterized the problem’s difficulty without a mechanism to run. Reproducing WaterDrum’s verification methodology as an explicit baseline shows it inherits a specific, previously-only-stated blind spot that SEAL-W’s independent diversity channel closes; five proven formal results (four theorems and a proposition) state precisely, and only, what is guaranteed; and the experimental evaluation reports where the certificate’s power degrades – at large per-client data scale, and at small calibration-sample size – as plainly as where it succeeds.

REFERENCES

- [1] European Parliament and Council of the European Union, “Regulation (eu) 2016/679 (general data protection regulation),” Official Journal of the European Union L119, 2016.

- [2] L. Tang, J. Joshi, and A. Kundu, "Characterizing privacy-audit alignment in behavioral audit of machine unlearning," *arXiv preprint arXiv:2606.14518*, 2026, v2.
- [3] T. L. Nguyen, M. Tuler de Oliveira, A. Braeken, A. Y. Ding, and Q.-V. Pham, "Towards verifiable federated unlearning: Framework, challenges, and the road ahead," *IEEE Internet Computing*, 2026, arXiv:2510.00833.
- [4] X. Lu, X. Niu, G. K. R. Lau, B. T. C. Nhung, R. H. L. Sim, J. R. Himawan, F. Wen, C.-S. Foo, S.-K. Ng, and B. K. H. Low, "Waterdrum: Watermarking for data-centric unlearning metric," in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2026, arXiv:2505.05064.
- [5] P. Maini, Z. Feng, A. Schwarzschild, Z. C. Lipton, and J. Z. Kolter, "Tofu: A task of fictitious unlearning for llms," in *Conference on Language Modeling (COLM)*, 2024.
- [6] M. Bun and T. Steinke, "Concentrated differential privacy: Simplifications, extensions, and lower bounds," in *Theory of Cryptography Conference (TCC)*, 2016, pp. 635–658.
- [7] P. Paillier, "Public-key cryptosystems based on composite degree residuosity classes," in *Advances in Cryptology—EUROCRYPT '99*. Springer, 1999, pp. 223–238.
- [8] J. Kirchenbauer, J. Geiping, Y. Wen, J. Katz, I. Miers, and T. Goldstein, "A watermark for large language models," in *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023.
- [9] A. Godinot, S. Azogagh, J. Ferry, and S. Gambis, "Manipulation-proof oblivious audits against deceptive model providers," *arXiv preprint arXiv:2608.04365*, 2026.
- [10] Y. Cao and J. Yang, "Towards making systems forget with machine unlearning," in *Proceedings of the 2015 IEEE Symposium on Security and Privacy (S&P)*, 2015, pp. 463–480.
- [11] L. Bourtoutle, V. Chandrasekaran, C. A. Choquette-Choo, H. Jia, A. Travers, B. Zhang, D. Lie, and N. Papernot, "Machine unlearning," in *Proceedings of the 2021 IEEE Symposium on Security and Privacy (S&P)*, 2021, pp. 141–159, arXiv:1912.03817.
- [12] C. Guo, T. Goldstein, A. Hannun, and L. van der Maaten, "Certified data removal from machine learning models," in *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020, arXiv:1911.03030.
- [13] A. Sekhari, J. Acharya, G. Kamath, and A. T. Suresh, "Remember what you want to forget: Algorithms for machine unlearning," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021, arXiv:2103.03279.
- [14] G. Liu, X. Ma, Y. Yang, C. Wang, and J. Liu, "Federaser: Enabling efficient client-level data removal from federated learning models," in *Proceedings of the 2021 IEEE/ACM 29th International Symposium on Quality of Service (IWQoS)*, 2021.
- [15] X. Gao, X. Ma, J. Wang, Y. Sun, B. Li, S. Ji, P. Cheng, and J. Chen, "Verifi: Towards verifiable federated unlearning," *IEEE Transactions on Dependable and Secure Computing*, 2024, arXiv:2205.12709; doi:10.1109/TDSC.2024.3382321.
- [16] D. M. Sommer, L. Song, S. Wagh, and P. Mittal, "Athena: Probabilistic verification of machine unlearning," *Proceedings on Privacy Enhancing Technologies (PoPETs)*, vol. 2022, no. 3, pp. 268–290, 2022, arXiv:2003.04247.
- [17] T. González, M. Dulce-Rubio, A. Ramdas, and M. Ribero, "Sequentially auditing differential privacy," *arXiv preprint arXiv:2509.07055*, 2025.
- [18] G. K. R. Lau, X. Niu, H. Dao, J. Chen, C.-S. Foo, and B. K. H. Low, "Waterfall: Framework for robust and scalable text watermarking and provenance for llms," in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024.
- [19] J. Dong, A. Roth, and W. J. Su, "Gaussian differential privacy," *Journal of the Royal Statistical Society Series B*, vol. 84, no. 1, pp. 3–37, 2022.
- [20] M. Jagielski, J. Ullman, and A. Oprea, "Auditing differentially private machine learning: How private is private sgd?" in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020, arXiv:2006.07709.
- [21] M. Nasr, S. Song, A. Thakurta, N. Papernot, and N. Carlini, "Adversary instantiation: Lower bounds for differentially private machine learning," in *Proceedings of the 2021 IEEE Symposium on Security and Privacy (S&P)*, 2021, pp. 866–882, arXiv:2101.04535.
- [22] P. W. Koh and P. Liang, "Understanding black-box predictions via influence functions," in *Proceedings of the 34th International Conference on Machine Learning (ICML)*, 2017.
- [23] M. S. Charikar, "Similarity estimation techniques from rounding algorithms," in *Proceedings of the 34th Annual ACM Symposium on Theory of Computing (STOC)*, 2002, pp. 380–388.
- [24] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. Agüera y Arcas, "Communication-efficient learning of deep networks from decentralized data," in *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2017.
- [25] Dua, D. and Graff, C., "Statlog (german credit data) data set," UCI Machine Learning Repository, 2017.
- [26] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, "Gradient-based learning applied to document recognition," *Proceedings of the IEEE*, vol. 86, no. 11, pp. 2278–2324, 1998.
- [27] J. A. Blackard, "Coverttype data set," UCI Machine Learning Repository, 1998.
- [28] R. Avenhaus, B. von Stengel, and S. Zamir, "Inspection games," in *Handbook of Game Theory with Economic Applications*, R. J. Aumann and S. Hart, Eds. North-Holland, 2002, vol. 3, ch. 51, pp. 1947–1987.
- [29] Court of Justice of the European Union, "Google spain sl and google inc. v agencia española de protección de datos (aepd) and mario costejá gonzález," Case C-131/12, Judgment of the Court (Grand Chamber), ECLI:EU:C:2014:317, 2014.